


ARCHITECTURAL PLAYBOOK

The AI Harness

Architecture, Authority, and Practice

Engineering a secure perimeter around non-deterministic inference engines.



Compiled by Douglas Mun with AI assistance.

Version 1.0 · July 2026

Table of Contents

Table of Contents.....	2
Executive Summary.....	3
Abstract.....	3
1. Origin of the Term.....	4
2. Two Senses of "Harness".....	4
3. What Is an AI (Agent) Harness?.....	5
4. Threat Model and Trust Assumptions.....	6
5. Why Harnesses Are Necessary.....	7
6. Core Components.....	7
6.1 The Agent Loop (foundational).....	8
6.2 Tool Design (the model–world interface).....	8
6.3 Context, Provenance & Memory (the scarce, taintable resource).....	11
6.4 Authority & Execution Control (what makes deployment safe).....	14
6.4.1 The authority model.....	14
6.4.2 Where checks sit.....	16
6.4.3 Execution safety: time, state, and repetition.....	16
6.4.4 Enforcement mechanisms — and their honest limits.....	17
6.4.5 Design tensions.....	18
6.5 Evaluation, Verification & Adversarial Testing (closing the loop).....	19
6.6 Operational Concerns (what production adds).....	20
7. A Worked Example.....	21
8. Harnesses and Capability Elicitation.....	23
9. Observability as Auditability.....	23
10. Architectural Patterns (secondary).....	24
11. The Security Surface (threat catalogue behind the §6.4 controls).....	24
12. Design Principles.....	26
13. Failure Modes.....	27
14. Minimum Adversarial Test Suite.....	28
15. The Shifting Model / Harness Boundary.....	29
16. Limitations & Open Problems.....	30
17. Harness vs. Model: A Clarifying Summary.....	32
18. Closing Note.....	33
19. References & Further Reading.....	34

Executive Summary

A large language model, on its own, produces plausible text—it cannot reliably take actions, remember what happened many steps ago, verify its own work, or stay within safe operating bounds. The **harness** is the engineering layer that wraps a model, converting raw generative capability into a dependable, tool-using, goal-directed system. The model sets the ceiling of what is possible; the harness determines how much of that ceiling is realized, and whether the power realized stays under the control of the person it serves.

This document defines the term, states an explicit threat model in which the adversary may control the request, retrieved content, tool results, peers, repetition, and external state—placing the model *inside* the blast radius rather than on the perimeter—and then concentrates on the four disciplines that most determine success: **tool design, context and memory, authority and execution control, and evaluation**. Around these it covers architecture patterns, the security threat catalogue, failure modes, a minimum adversarial test suite, and the shifting boundary between model and harness.

Two themes recur. First, *tool availability is never authorization*: every consequential action must be authorized per invocation against concrete values, fresh state, and the principal's original mandate. Second, protection observed today is partly *borrowed*—a more capable model can complete exploits that a weaker one merely began, so every model upgrade is a security-relevant change. Prompt injection remains unsolved; the engineering goal is to mitigate and contain it, and to prove—against an adaptive adversary—that the containment holds.

Abstract

Large language models (LLMs) are powerful but fundamentally unstructured: given text, they produce text. On their own, they cannot reliably take actions, remember what happened many steps ago, verify their own work, or stay within safe operating bounds. The **harness** is the engineering layer that wraps a model, turning raw generative capability into a dependable, tool-using, goal-directed system. This document defines the term (including its two senses in machine learning), states an explicit threat model, and then concentrates on the four elements that most determine whether a harness succeeds — **tool design, context and memory, authority and execution control, and evaluation** — before covering architecture, the security surface, failure modes, adversarial testing, and the evolving boundary between model and harness.

How to read this document. Not all components matter equally. Sections §6.2–§6.5 (tools, context/memory, authority & execution control, evaluation) are the load-bearing ones and are treated at length. This document is an architectural guide with security engineering woven in; it is not, by itself, a complete secure-implementation specification — §16 states what remains open.

1. Origin of the Term

The word *harness* comes from the Old French *harnois*, originally meaning military equipment, later narrowing to the straps and fittings that hitch a draft animal to a load. The engineering idea carried over unchanged: **take a powerful but undirected source of energy and channel it into controlled, useful work.**

In machine learning, the word arrived first through **evaluation**, not agents. Researchers wrote *test* and *evaluation harnesses* — scaffolding that feeds a model a benchmark, collects its outputs, and scores them — well before "AI agents" were common. Only later did the same word come to be used to describe the runtime scaffolding around a deployed, tool-using model. Both senses are current.

2. Two Senses of "Harness"

	Evaluation harness	Agent harness
Purpose	Measure capability on a fixed benchmark	Accomplish real tasks at runtime
Runs	Offline, batched, reproducibly	Online, interactively, per request
Examples	lm-evaluation-harness , SWE-bench's harness	A coding assistant's loop, a support agent
Key concern	Fair, reproducible scoring	Reliability, safety, cost, latency

The two connect through §8: because a harness can dramatically change measured performance, a benchmark score is only meaningful relative to the harness that produced it. This document focuses on the **agent harness**, returning to evaluation where they intersect.

3. What Is an AI (Agent) Harness?




Figure 1. The engine vs. the drivetrain — raw model output versus directed, verifiable work.

An **agent harness** (also called an *agent loop* or *scaffold*) is the software system that constructs the model's input, executes the actions it requests, feeds the results back, enforces constraints, manages state and authority across turns, and decides when the task is done.

On the model's role, precision matters. At the application-security boundary, **each model invocation should be treated as lacking durable, trusted task state unless that state is explicitly supplied and authenticated by the surrounding system.** Some provider APIs do preserve conversational or execution state between calls; the point is not literal statelessness but ownership — durable memory, authorization, and auditability are harness responsibilities, never the model's. Everything that makes an LLM behave like an *agent* — something that perceives, decides, and acts on someone's behalf over time — lives in the harness, not the weights.

Analogy: A horse harness converts raw pulling strength into directed work — a straight furrow instead of a bolt for the horizon. An agent harness converts raw predictive capability into directed work — a completed task instead of plausible-sounding text.

4. Threat Model and Trust Assumptions




Figure 2. The threat model — mapping the attack surface around the agentic system.

Any serious harness design begins by stating what the adversary may control. Assume an attacker may control one or more of:

1. **Direct user input** — the request itself may be adversarial.
2. **Retrieved content** — a webpage, email, document, repository, database row, or search result the agent reads.
3. **Tool results and tool metadata** — outputs, error messages, and even the descriptions and schemas of tools.
4. **Third-party components** — a plugin, MCP server, connector, or peer agent in a multi-agent system.
5. **Repetition** — repeated requests that exploit stochastic behavior until an attack lands.
6. **External state** — the world may change between when an action is authorized and when it executes.
7. **The approval surface** — the UI or channel through which a human reviews and approves actions.

Attacker objectives include unauthorized actions, data exfiltration, **persistence through memory**, privilege escalation, denial of service (including cost exhaustion), audit evasion, and inducing the agent to act outside the principal's actual intent.

Two trust assumptions follow and recur throughout this document:

- **The model is inside the blast radius, not on the perimeter.** Once untrusted content enters the model's context, the model may follow it. Controls that depend on the model resisting manipulation are mitigations, not boundaries.
- **Tool availability is not authorization.** That a mail tool is exposed does not authorize sending any recipient any attachment; that a payment tool exists does not authorize every amount and destination. Authorization must be evaluated per invocation, against concrete values (§6.4).

5. Why Harnesses Are Necessary

Limitation of the raw model	What the harness provides
No durable, trusted task state	Context, memory, and state management
Cannot execute code, browse, or touch files	Tool / function-calling infrastructure
Cannot independently verify its own outputs	External tests, authoritative data, verification
No concept of "stop when done"	Termination conditions, step and cost limits
Follows instructions wherever they appear	Provenance tracking, authority controls, human review
Non-deterministic, sometimes wrong	Retries (with idempotency), self-correction, error handling

6. Core Components




Figure 3. The architectural blueprint — the four load-bearing pillars of a defensible agent loop.

A harness has several parts, but they do not carry equal weight. The **agent loop** (§6.1) is foundational but simple. The four components that determine how good — and how safe — an agent is are **tool design (§6.2)**, **context & memory (§6.3)**, **authority & execution control (§6.4)**, and **evaluation (§6.5)**. A caution on vocabulary: the popular umbrella term "guardrails" bundles together several distinct engineering disciplines — authorization, execution safety, approval integrity, and enforcement — which §6.4 treats separately, because conflating them is itself a source of design error. The subsection lengths deliberately reflect this ordering of importance.

6.1 The Agent Loop (foundational)

1. Assemble context (system prompt + history + prior tool results, with provenance)
2. Call the model
3. Parse the response strictly
4. If it proposes an action → authorize → (approve if required) → execute → verify → append result → go to 1
5. If it returns a final answer → validate → return it
6. If a limit is hit (steps, time, cost, tokens) → stop and report

This control loop is often *implemented* with the **ReAct** pattern (Reason + Act; Yao et al., 2022), where the model interleaves explicit reasoning with actions. Be precise, though: ReAct is one prompting strategy for driving the loop, not a synonym for it. Note also the vocabulary in step 4: the model *proposes*; the harness *authorizes*, *executes*, and *verifies*. The loop itself is simple; the difficulty — and the value — lives in the components that follow.

6.2 Tool Design (the model–world interface)




Figure 4. Pillar 1 — tool design and the supply chain.




Figure 5. The mid-session supply-chain vulnerability.

Tools give the model "hands," and **tool design is often one of the highest-leverage surfaces in a harness** — in many deployed systems, *the* highest — though identity architecture, data quality, or isolation can dominate in others. Two harnesses wrapping the same model, differing only in how their tools are named, scoped, and documented, can differ enormously in reliability. The reason is fundamental: **the model acts almost entirely on a tool's description, not its implementation.** A tool the model misunderstands is a tool it misuses.

What the harness owns for every tool call:

- **Validation** — check arguments against a schema *before* executing. Necessary, but understand its limits: schema validation proves an argument is *syntactically* valid, not that the action is authorized or safe (§6.4 owns that).
- **Execution** — perform the underlying action under a scoped capability (§6.4), never ambient credentials.
- **Result formatting** — return something the model can use: shaped for interpretation, truncated when large, and **carrying provenance** — the result of a tool that touched untrusted content is itself untrusted (§6.3).
- **Failure handling** — surface timeouts and errors back to the model as readable text so that it can adapt rather than reason from silence. But treat error text as a channel: an attacker-influenced tool can embed instructions in its *error messages*, so errors are data too.

What makes a tool set good:

- **Descriptions are prompt engineering — and security-sensitive configuration.**
The name, parameter docs, and usage notes are the entire interface the model sees. Ambiguous descriptions cause misuse; *malicious* descriptions cause hijacking. Write them as carefully as any prompt, and control changes to them as carefully as code (see supply chain below).
- **Right granularity.** Too coarse, and one call does too much and hides failures; too fine, and simple tasks need long, error-prone chains. Aim for actions a competent operator would recognize as a single step.
- **Few, non-overlapping tools.** Selection accuracy degrades as the menu grows and blurs. A small, orthogonal set outperforms a large, redundant one.
- **Legible errors.** A tool that fails with a clear, actionable message lets the model recover; one that fails opaquely sends it into loops.
- **Side-effect awareness.** Read-only tools are safe to retry; mutating tools are not, and need idempotency keys (§6.4). The harness must know which is which.

Tool supply chain. Because descriptions and schemas steer the model, they are part of the attack surface. An adversary may alter a description after review, publish a malicious MCP server or plugin, change a schema while keeping the name, return deceptive metadata, compromise a dependency, or "rug-pull" after installation. OWASP's agentic-security guidance lists supply-chain compromise and tool misuse among the top agentic risks. Minimum controls: **signed and pinned tool manifests; version and schema hashes in audit records; allowlisted publishers; capability re-review whenever a description or schema changes; separation of tool discovery from tool authorization; and egress/credential policies that do not depend on the tool's own claims.**

The dynamic-definition blind spot. Pinning assumes tool definitions are static artifacts — but protocols like MCP deliberately move tool definition to a live server, which may return schemas and descriptions *dynamically at runtime*. That reopens the supply chain mid-session: an allowlisted, individually secure parsing server whose returned metadata varies with the document it processes can be driven by an attacker-controlled document, to present an altered schema or description — converting a safe tool into a confused-deputy vector *between* manifest reviews. The pinning discipline must therefore extend to runtime: **freeze the tool surface at session start** (schemas and descriptions captured and hashed at grant creation); **treat any mid-session change to a tool's definition as untrusted input**, not as configuration — diff it, quarantine the tool, and require re-review before the changed definition reaches the model; and never allow content a tool has *processed* to feed back into that tool's own advertised capabilities within the session.

Most "the agent is unreliable" problems, on inspection, turn out to be tool problems — a description the model misread, an error it couldn't see, a menu too crowded to choose from. Fixing tools is usually cheap and high-impact. But tool *trust* problems are security problems, and they need supply-chain discipline, not just better prose.

6.3 Context, Provenance & Memory (the scarce, taintable resource)




Figure 6. Pillar 2 — context budget and the taint tracker.




Figure 7. The threat of memory laundering.

The model's context window is a **finite, shared budget**, and on any single call, the model sees one flat blob of tokens assembled by the harness. Managing that blob — what to include, when,

in what form, and *from where* — is the second core discipline. Two things go wrong: the task runs out of room or drowns in irrelevance (a capacity problem), and untrusted content gains influence it should never have (a provenance problem). Both are the harness's job.

Core strategies for a growing task:

- **Sliding window** — keep only the most recent N turns.
- Summarization/compaction — periodically compress older history.
- **Retrieval-augmented context** — store information externally; pull in only what the current step needs.
- **Structured scratchpads** — persistent notes or files the model reads and writes across a long task.

Loading discipline. The strategies reduce to one question asked repeatedly: *what earns a place in the window this turn?* Concrete mechanisms (using names common in coding assistants like Claude Code as examples of general patterns) differ mainly in when they load:

- **Always-on context** — standing instructions and project configuration loaded at session start and kept throughout (e.g. a `CLAUDE.md` project file). Never fetched, always present; must stay lean because it occupies budget every turn.
- **On-demand / progressive disclosure** — only a cheap *pointer* is always available; the full payload loads when relevant. Skill files (`SKILL.md`) are the clearest example: the short description is always in context, the body is read only when the task calls for it. Loading everything every turn would blow the budget.
- **Cross-session memory** — state that outlives a conversation, written now and retrieved later.
- **Accumulating context** — history and tool results grow turn over turn and are periodically compacted.

Provenance is not optional. Every piece of content entering the window should carry metadata: where it came from, what trust level that source has, and what data classification it holds — and those labels must be **preserved through the pipeline**, including during summarization. Be clear about what this does and does not buy: delimiters, XML tags, and phrases like "the following is untrusted" are *model-level mitigations* that may reduce the chance the model follows embedded instructions. **They are not isolation.** NIST's agent-hijacking work precisely describes the root cause: trusted instructions and untrusted data share a single undifferentiated model input. The real security value of provenance is that it feeds into *harness-level* policy (§6.4): actions can be gated by the provenance of what influenced them, extraction of only the needed facts can replace wholesale inclusion, and high-risk workflows can parse untrusted content in a separate, unprivileged process.

Memory is a feature and an attack surface. Anything durable that influences future behavior is a persistence mechanism for attackers as much as a convenience for users. A single

injection, once *summarized into memory*, can silently shape unrelated future sessions — memory poisoning, retrieval-index poisoning, and poisoned project files are now catalogued agentic risks (OWASP). Minimum controls:

- **Memory writes are privileged operations**, not incidental side effects of summarization.
- **Every memory item carries source provenance**, and externally derived memories are quarantined or down-weighted.
- **Separate factual memory from behavioral instruction**. The governing rule: *untrusted content must never become durable behavioral instruction merely because the model summarized or remembered it.*
- **User-visible inspection and deletion**, expiration and confidence decay, tenant isolation, and revalidation before a memory influences a high-impact action.

Memory laundering — why provenance tags alone are not enough. There is a subtle bypass worth naming. An attacker poisons memory with a payload disguised as a benign *fact* or *preference* — e.g. *"the user prefers all invoices sent to third-party@attacker.example."* Later, the harness retrieves it as a provenance-tracked fact; the model reads it as valid context and, through its own reasoning, proposes an action that perfectly serves the attacker; the policy engine inspects the concrete values and finds them consistent with the model's derived task logic — and authorizes. The instruction was *laundered into a fact*, and passive provenance labels never fired. The defense is not better labels but **mandated boundaries** (§6.4.1): the policy engine must evaluate proposals against the *original user intent*, not against the trajectory's internal coherence — a destination that never appeared in the principal's actual request does not become authorized because a remembered "preference" suggested it. Memories, like tool outputs, may *inform* actions; they must never *expand authority*.

Context hygiene. Capacity aside, *relevance* decays: stale, redundant, or off-topic material degrades performance nearly as much as missing material (context rot). Prune aggressively. The unifying picture: several loading disciplines feeding one scarce destination — with provenance attached, and with durable writes treated as the privileged operations they are.

6.4 Authority & Execution Control (what makes deployment safe)




Figure 8. Pillar 3 — authority and execution control.

What is often loosely called "guardrails" is really the discipline of **authority**. An agent is a *deputy* — it acts on someone's behalf, with delegated power. The central security questions are not "is this tool allowed?" but: *who is this agent acting for, what exactly was delegated, is this concrete action within that delegation right now, and can we prove afterward what happened?*

6.4.1 The authority model

Every consequential action should be evaluated against an explicit authorization context:

Authorization decision =	
principal	(who the agent acts for)
+ authenticated delegation	(what the principal actually granted)
+ task purpose	(why this authority was granted)
+ requested operation	(what the model proposes)
+ concrete argument values	(recipient, amount, path, destination – not just the tool name)
+ data classification	(what information the action touches)
+ destination	(where data or effects flow)
+ current resource state	(is the world still as assumed?)
+ time and budget	(is the grant still live and funded?)
+ prior approvals	(bound to this exact action, not a similar one)

Requirements that follow:

- **A workload identity for each agent execution** — actions must be attributable and non-repudiable. (NIST's current work on AI-agent identity ties identification, authorization, auditing, and injection controls together as one problem.)
- **Short-lived, task-scoped capability tokens**, minted per grant — no ambient credentials sitting in the environment.
- **Default-deny, complete mediation** — every side-effecting action passes the policy engine; nothing is implicitly allowed.
- **Value-level authorization** — the decision considers concrete arguments, not merely which tool was invoked. A schema-valid transfer to an unauthorized account must be denied *on the values*.
- **Explicit delegation chains for subagents** — a spawned agent holds a *narrowed* grant, never the parent's full authority, and peer-agent messages confer no authority at all.

The mandate, not just the scopes. Traditional scopes and role-based access control answer "*may* this identity call this API?" — necessary but insufficient for an autonomous system, because an agent can assemble individually permitted actions into an unauthorized whole (the memory-laundering case in §6.3 is one instance). The grant should therefore carry a **mandate**: a digest of the principal's actual intent plus continuous-validation boundaries (budget ceilings, allowed destination and domain tiers, data-classification limits). The harness then checks not only each action but the **trajectory**: if the agent dynamically acquires tools, aggregates data across domains, or targets destinations in a way that shifts the task's risk tier away from the original mandate, that structural drift itself triggers an alert or forces human review — even when every individual call would pass value-level policy. This is continuous authorization, not one-time consent.

Mapping to enterprise infrastructure. These requirements need not be bespoke; they map onto existing zero-trust standards, which a production implementation should reach for rather than reinvent: **workload attestation** frameworks (e.g. SPIFFE/SPIRE) for issuing short-lived, cryptographically verifiable identities to agent loops and subagents; **OAuth 2.x token exchange (RFC 8693)** for modeling downstream delegation, where an agent exchanges its primary token for a down-scoped, short-lived token before calling a tool or spawning a subagent; and attribute-based policy frameworks (e.g. NGAC-style models) for expressing value-level decisions. The paper's authority model is the *specification*; these are candidate *implementations*.

6.4.2 Where checks sit




Figure 9. Checkpoints of an agentic action.

- **Input** (before the model call) — attach provenance and classification; strip secrets.
- **Output** (after the response, before acting) — strict parsing; policy screening of the proposed action.
- **Pre-execution** — the authorization decision above, evaluated against **fresh authoritative state**, immediately before commit.
- **Post-execution** — verify postconditions: did the action do what was authorized, and only that? A compromised tool can return forged success; verification against independent state catches it.

6.4.3 Execution safety: time, state, and repetition

Schema validation proves syntax; it says nothing about *when* the action runs. The gap between authorization and execution is an attack surface (time-of-check-to-time-of-use):

- A file swapped after review but before upload; a payment destination changed after approval; permission revoked mid-task; inventory changed between plan and purchase.
- **Controls:** precondition versions or entity tags; re-authorization immediately before commit; optimistic concurrency or locks; approvals carrying nonces and expiry.

Repetition is the other gap. A timeout followed by a retry can execute a payment twice; two parallel subagents can issue conflicting mutations:

- **Controls:** idempotency keys on every mutating call; replay protection; separate retry policies for reads (retry freely) and writes (retry only with idempotency); serialized

mutations by default, with parallelism reserved for demonstrably independent read-only work; transaction boundaries and compensating actions for multi-step effects; explicit partial-failure handling and cancellation propagation.

Stochastic action flooding. Concurrency controls are themselves attackable. Because agents are stochastic and can spawn parallel subagents, an injection can induce a burst of near-identical, near-simultaneous mutation proposals. Against baseline optimistic concurrency, the flood produces continuous conflict-rollback churn — a state-drift denial-of-service — and, worse, pressure on exactly the paths most likely to be mishandled: timeouts in the policy engine or audit logger that an implementation quietly treats as fail-open. **Controls:** rate-limit *proposals* (not just executions) per grant; deduplicate semantically-equivalent pending mutations before they reach the policy engine; apply backpressure — a grant whose conflict-rollback rate spikes is suspended for review, not retried harder; and re-verify that every timeout path on the control plane fails closed (§6.4.5), because a flood is precisely how an attacker goes looking for the one that doesn't.

6.4.4 Enforcement mechanisms — and their honest limits




Figure 10. Enforcement mechanisms — and their honest limits.

Three mechanisms, with different trust properties:

- 1. Deterministic/programmatic** — policy engines, allowlists, validation, in-harness code. An attacker cannot *persuade* a correct policy engine the way they can persuade a model. But deterministic enforcement is **necessary, not sufficient**: it must be complete, fail-closed, value-aware, and placed at the execution boundary. Attackers do not argue with **if** statements; they exploit missing checks, wrong policy, canonicalization and parser

differentials, stale state, downstream injection (model output reaching a shell, SQL, or unescaped template), and fail-open error paths.

2. **Model-based** — a classifier or second LLM judging the action. Flexible for fuzzy calls; probabilistic and itself manipulable. Layers on top of deterministic checks, never replaces them.
3. **Human-in-the-loop** — strong **only under specific conditions**. A human can be shown a misleading summary, a truncated recipient list, a benign preview followed by altered parameters, or an approval screen containing attacker-controlled text — and too many prompts produce habituated click-through. Meaningful approval requires: the **exact** effects displayed (recipients, amounts, files, permissions); changes from the user's original request highlighted; attacker-controlled text excluded from trusted UI chrome; the approval **cryptographically bound to a canonical digest of the exact action**; short expiry; no replay; re-approval after any material change; and independent channels for the highest-impact operations. Stated precisely: *human approval provides strong control only when it is informed, resistant to fatigue and spoofing, and bound to the exact action and state being authorized.*

6.4.5 Design tensions

- **Fail-closed vs fail-open** — when a check *itself* errors (including the audit logger being unavailable), high-impact actions block.
- **False positives vs false negatives** — over-blocking strangles work; under-blocking admits harm. Tune per impact tier.
- **Friction erodes control** — reserve human approval for actions that warrant it, or users will click through everything.
- **Defense in depth** — no single check is load-bearing: identity *and* value-level policy *and* isolation *and* egress limits *and* postcondition verification.

Impact tiering ties it together: auto-approve low-impact, reversible actions under a valid grant; validate and re-authorize medium ones; bind high-impact or irreversible ones to exact-action human approval. Proportionate, not uniform.

6.5 Evaluation, Verification & Adversarial Testing (closing the loop)




Figure 11. Pillar 4 — verification and adversarial testing.

Capable harnesses verify work — but be precise about what models can and cannot do here.

Models can perform useful self-review, but self-review is correlated with the original generation and does not constitute independent verification.

High-confidence verification requires external evidence: tests that execute, authoritative data, formal constraints, or independently designed checks. Iterative self-correction — act, observe the failure, revise — is the mechanism behind approaches such as Reflexion (Shinn et al., 2023), and it works best when the "observe" step is grounded in something external to the model.

Capability evaluation:

- **Outcome vs. trajectory.** Outcome eval asks whether the result is correct; trajectory eval inspects the path — tools called, steps wasted, recovery from errors. Outcome metrics are cheaper and less gameable; trajectory metrics diagnose *why* failures happen.
- **Reliability over runs.** Agents are stochastic; a single run is noise. Report success *rates* over many trials. `pass@k` (success within k attempts) is informative but **not inherently desirable**: retrying amplifies benign success and adversarial success alike, so pair it with per-attempt reliability and cost-normalized figures.
- **LLM-as-judge, with caveats.** Scales well; inherits the grader's biases and is itself manipulable. Sits alongside deterministic checks, not in place of them.

Adversarial evaluation is a separate discipline, and capability metrics do not cover it.

NIST/CAISI's agent-hijacking evaluations, built on AgentDojo, found that attacks adapted to the target system dramatically outperform fixed baselines — novel red-teamed techniques reached

81% hijacking success where the strongest known baseline attacks achieved 11% — and that testing across multiple attempts yields more realistic results than single-shot evaluation. Fixed public attack sets give false comfort. A security evaluation needs: adaptive attackers and held-out attack families; repeated attempts; compromised-tool and compromised-peer scenarios; persistence and recovery testing; and measurement against the *exact* production policy and tool configuration.

Key metrics:

Metric	Purpose
Clean task success rate	Retained utility alongside safety
Unauthorized-action rate	Actual policy violations
Attack success rate	Attacker objective completion
Partial-compromise rate	Initial obedience even when the full exploit fails
Impact-weighted risk	A leaked benign email ≠ exfiltrated credentials
False-deny rate	Excessive friction
Repeated-attempt compromise	Stochastic amplification
Maximum blast radius	Containment
Detection & containment time; recovery correctness; audit completeness	Operational response and forensics

6.6 Operational Concerns (what production adds)

- **Cost.** Every loop iteration is one or more model calls; long trajectories and large contexts compound. Budget checks belong *before* the spend is incurred where possible, and cost caps serve both as budget control and as a denial-of-service defense — per-principal quotas and circuit breakers contain an attacker who drives extreme token or API usage.
- **Latency.** Agentic tasks are sequential; end-to-end latency is the *sum* of steps. Stream partial output, parallelize only independent read-only calls, minimize round-trips.
- **Reproducibility.** Model and environmental nondeterminism usually prevent *exact* replay. The achievable standard is **forensic reconstruction**: with full trajectories logged, versions pinned, and decisions recorded, an investigator can establish what happened and why, even if the run cannot be repeated bit-for-bit.

7. A Worked Example




Figure 12. System in motion — naive vs. defensible lifecycle.

The naive loop — call the model, execute whatever it asks, append the result — is pedagogically tempting and **unsafe as an implementation pattern**: it pipes tool output straight back into context (an injection channel), authorizes at the tool level only, checks cost after spending, and cannot bind an approval to an action. Here is the shape of a defensible lifecycle instead:

```

# 1. Establish authority before any model call
grant = create_task_grant(
    principal=user_identity,      # who this agent acts for (§6.4.1)
    purpose=normalized_intent,    # why authority was granted
    scopes=minimum_scopes,       # least privilege
    expires_at=deadline,         # tokens, dollars, steps
    budgets=budgets,
)
history = [system_prompt(), user_message(task)]  # trusted roots

while within_limits(grant):      # pre-spend budget check
    context = build_context(history, provenance=True)  # §6.3 labels preserved

    proposal = parse_strictly(
        model.propose(context, tool_schemas(grant))
    )  # §6.1 model *proposes*

    if proposal.is_final:
        return validate_and_release(proposal, grant)  # §6.5 output validation

    state = fetch_authoritative_state(proposal)  # §6.4.3 fresh state

    decision = policy.evaluate(
        grant=grant, action=proposal,
    )  # §6.4.1 value-level, # default-deny,
  
```

```

        args=proposal.args, state=state,                # provenance-aware
        provenance=proposal.input_provenance,
    )
    if not decision.allowed:
        history.append(denial_record(proposal, decision)); continue

    if decision.requires_approval:                      # §6.4.4 bound approval
        approve_exact_action(digest=canonical_digest(proposal),
                             diff=render_effects(proposal))

    raw = execute(                                       # §6.4.3 safe execution
        proposal,
        capability=decision.short_lived_capability,     # no ambient creds
        idempotency_key=stable_action_id(proposal),     # no duplicate effects
        sandbox=decision.sandbox_policy,
        egress=decision.egress_policy,
    )

    result = assess_taint(raw, proposal.input_provenance) # §6.3 taint computed
                                                         # BEFORE any commit
    if verify_postconditions(result, proposal, state):   # §6.4.2 don't trust
        commit_state_change(result)                    # forged success
        history.append(result_with_provenance(result))
    else:
        trigger_compensating_transaction(proposal)      # undo, don't shrug
        history.append(denial_record(proposal,
                                     "postcondition failure"))

    audit.write_tamper_evident_record(grant, proposal,   # §9 forensics
                                     decision, result)

    history = compact_if_needed(history)                # §6.3 memory
return stop_with_redacted_diagnostics(grant)           # no secret leakage

```

Mutating calls are serialized; parallelism is reserved for independent read-only work. Note the ordering discipline in the tail of the loop: taint is assessed on the *raw* result before anything mutates system state, commit happens only after postconditions pass, and a failed verification produces a compensating transaction plus an explicit denial record — never a silent shrug. The model still appears on one line — but note what surrounds it now: identity, delegation, fresh-state authorization on concrete values, bound approval, scoped ephemeral capability, idempotency, taint-then-commit ordering, postcondition verification, and a tamper-evident record. That surround is the harness.

8. Harnesses and Capability Elicitation

A central idea for building or evaluating frontier systems: **the same model performs very differently under different harnesses**. Harness quality is a form of *capability elicitation* — how much of a model's latent ability is actually surfaced.

1. **Benchmark scores are harness-relative.** "Model scores X% on Y" is incomplete without the harness. Better tools, cleaner context, retries, or self-verification move scores substantially with no change to the weights. Fair model comparison holds the harness fixed.
2. **Harness improvement is often the cheapest lever.** When a deployed agent underperforms, the bottleneck is usually scaffolding — tool ergonomics (§6.2), lossy context (§6.3), missing error feedback — not the model.

It cuts both ways for safety: a capability a weak harness fails to surface may be reachable by a stronger one, so conservative capability assessments assume a competent harness — and *adversarial* evaluations (§6.5) must assume a competent attacker.

9. Observability as Auditability

"Log everything" is not a security property; logs are themselves sensitive and attackable. They may contain access tokens, personal data, retrieved malicious prompts, generated exploit code, and system instructions — and an attacker who can evade or corrupt the audit trail can erase the incident. Requirements for **security-grade auditability**:

- **Structured events, not prose** — with correlation IDs linking model call → policy decision → approval → execution → verification.
- **Identity in every record** — principal, delegation chain, workload identity, tool manifest hash, model and policy versions.
- **Tamper-evident, append-only storage**, with access controls and separation of duties (the agent's credentials cannot rewrite its own audit trail).
- **Fail-closed coupling** — if the audit path is unavailable, high-impact execution blocks.
- **Redaction and tokenization of secrets**; retention and deletion policies; alerting on anomalous sequences; replay tooling over redacted snapshots.
- **Incident readiness** — containment procedures, credential revocation, and (per §6.3) rollback of poisoned memory.

10. Architectural Patterns (secondary)

Single-agent loop. One model, one tool set, one loop. Simplest and most common.

Planner–executor. A planning step decomposes a goal; an executor loop carries out subtasks. Suits long, multi-stage work.

Multi-agent / orchestrator. Specialized roles under an orchestrator. Robustness through cross-checking, at the cost of complexity, latency, and tokens — and a **trust cost**: inter-agent messages are untrusted input, peer instructions confer no authority, and each hand-off narrows (never widens) the delegation chain (§6.4.1). Multi-agent designs are frequently reached for prematurely; they pay off mainly when subtasks are genuinely parallelizable or benefit from independent perspectives. When subtasks are sequential and interdependent, a single well-designed loop is simpler and more reliable — every hand-off is a place for context and authority discipline, to be lost.

Human-in-the-loop. The harness pauses at checkpoints for review — subject to the approval-integrity requirements of §6.4.4.

11. The Security Surface (threat catalogue behind the §6.4 controls)

Where §6.4 covers the *controls*, this section catalogues the *threats*. The four classics first:

- **Prompt injection**, including the *indirect* variant where the instruction is planted in content the agent later retrieves (Greshake et al., 2023). No fully reliable defense exists; realistic web-agent benchmarks find low-effort injections cause agents to begin following adversarial instructions at high rates (WASP, 2025).
- **Confused deputy** (Hardy, 1988, applied to agents): an agent holding the user's authority is tricked into misusing it. The §6.4.1 authority model is the direct response.
- **Data exfiltration** via any outbound-capable tool; restricted egress and destination-aware data-loss controls respond.
- **Unsafe code execution**; isolation and no ambient credentials respond.

The fuller agentic threat catalogue (aligned with OWASP's agentic top-10 and NIST's hijacking work):

Threat	Essential control
Goal/plan hijacking	Revalidate objective and authority at each consequential transition
Memory & context poisoning	Privileged memory writes, provenance, quarantine, expiry, rollback (§6.3)
Memory laundering (instruction disguised as fact)	Mandate-boundary policy: memories inform, never expand authority (§6.3, §6.4.1)
Identity & privilege abuse	Workload identities, scoped delegation, no ambient authority (§6.4.1)
Tool / MCP supply-chain compromise	Signing, pinning, manifest hashes, re-review on change (§6.2)
Dynamic tool-definition mutation	Freeze tool surface at session start; mid-session changes quarantined as untrusted input (§6.2)
Stochastic action flooding	Proposal rate limits, dedup, backpressure, fail-closed timeouts (§6.4.3)
Improper output handling	Treat generated strings as hostile before shell, SQL, HTML, templates, APIs
Schema-valid semantic abuse	Value-level and business-policy authorization (§6.4.1)
Insecure inter-agent communication	Authenticate peers; peer content confers no authority (§10)
Replay & duplicate effects	Idempotency keys, nonces, approval expiry (§6.4.3)
Approval spoofing	Trusted UI chrome, exact-action digest binding (§6.4.4)
TOCTOU / state drift	Fresh-state authorization at commit, entity tags (§6.4.3)
Cascading failure	Circuit breakers, dependency budgets, isolation and rollback
Resource exhaustion	Per-principal quotas, concurrency limits, downstream API budgets (§6.6)
Audit evasion	Tamper-evident logs, independent telemetry, fail-closed coupling (§9)
Model/tool version drift	Pinning, canary evaluation, rollback
Secret exposure	Ephemeral credentials, destination-aware DLP, log redaction (§9)
Sandbox escape & side channels	Defense-in-depth isolation, monitored egress

The through-line: **untrusted input plus real authority equals attack surface**, and the harness — not the model — is where it is defended.

12. Design Principles



Figure 13. The architectural manifesto — twelve principles.

- 1. Authority is explicit** — every action is attributable to a principal via an authenticated, scoped, expiring delegation; tool availability is never authorization.
- 2. Least privilege, ephemerally held** — minimum scopes, short-lived capabilities, no ambient credentials.
- 3. Complete mediation, default-deny** — every side-effecting action passes policy, evaluated on concrete values and fresh state.
- 4. Provenance everywhere** — all content carries source and classification labels, preserved through the pipeline; policy acts on provenance, not the model's interpretation of it.
- 5. Untrusted content never becomes durable instruction** — memory writes are privileged; summarization is not a laundering step.
- 6. Approvals bind to exact actions** — digest-bound, expiring, non-replayable, faithfully rendered.
- 7. Effects are exactly-once and verified** — idempotency on writes, postcondition checks, compensation for partial failure.
- 8. Fail closed on the control path** — including when the audit trail is unavailable.
- 9. Observability is auditability** — structured, identity-bearing, tamper-evident, redacted (§9).
- 10. Bounded autonomy** — explicit stop conditions for steps, time, cost, and completion.
- 11. Errors are information — and data** — surface failures to the model legibly; treat their content as untrusted.
- 12. Context hygiene** — keep the window relevant; prune aggressively.

13. Failure Modes

Failure mode	Description	Typical mitigation
Tool misuse	Malformed or unsafe parameters	Schema validation; clearer descriptions (§6.2)
Context overflow	Task outgrows the window	Summarization, retrieval, scratchpads (§6.3)
Context rot	Stale/irrelevant material degrades output	Aggressive pruning (§6.3)
Loop divergence	Repeats actions or stalls	Step limits, progress checks, re-grounding
Unauthorized action	Effect outside the actual delegation	Value-level, fresh-state authorization (§6.4)
Duplicate effects	Retry or race causes double execution	Idempotency keys, serialized writes (§6.4.3)
Silent errors	Failures not surfaced; model reasons from bad state	Return errors as readable (untrusted) text
Forged success	Compromised tool reports completion falsely	Postcondition verification (§6.4.2)
Goal drift	Objective lost over a long task	Periodic re-statement; re-validate authority
Prompt injection	Untrusted content redirects the agent	Provenance-gated policy, egress limits, approval binding (§11)
Memory poisoning	Injected content persists across sessions	Privileged writes, quarantine, rollback (§6.3)

14. Minimum Adversarial Test Suite

The Minimum Adversarial Test Suite	
Demonstrate these – with adaptive attacks and repeated attempts, not fixed public sets.	
Test	Required Outcome
Injection in a retrieved email or webpage	No unapproved side effect
Poisoned cross-session memory	Quarantined, inspectable, removable
Malicious tool-description update	Tool rejected pending re-review
Schema-valid transfer to a bad account	Denied on the concrete values
Compromised tool returns forged success	Postcondition verification detects it
Audit logger unavailable	High-impact execution fails closed

Figure 14. The minimum adversarial test suite.

Before describing a harness as secure, demonstrate at least these outcomes — with *adaptive* attacks, not only fixed public sets, and repeated attempts per scenario (NIST's findings on adaptation and repetition make single-shot, fixed-set testing misleading):

Test	Required outcome
Injection in a retrieved email or webpage	No unapproved side effect
Injection hidden in a tool error message	Error handled as untrusted data
Poisoned cross-session memory	Quarantined, inspectable, removable
Malicious tool-description update	Tool rejected pending re-review
Schema-valid transfer to an unauthorized account	Denied on concrete values
Permission revoked during execution	Commit blocked at fresh-state check
Timeout followed by retry	Side effect occurs at most once
Two agents issue conflicting mutations	Conflict detected and contained
Approval preview differs from execution	Digest mismatch blocks execution
Peer agent supplies malicious instructions	No authority conferred
Encoded command reaches shell/SQL	Downstream validation blocks it
Compromised tool returns forged success	Postcondition verification detects it
Audit logger unavailable	High-impact execution fails closed

Attacker drives extreme token/API use	Quotas and circuit breakers contain cost
Repeated stochastic attack (25+ attempts)	Risk remains within defined threshold
Malicious instruction laundered into memory as a "fact"	Mandate check denies the derived action
MCP server mutates schema mid-session	Change quarantined; tool suspended pending re-review
Injection-induced flood of parallel mutations	Rate limits and backpressure contain it; all timeout paths fail closed
Model upgraded, harness unchanged	Full suite re-run pre-rollout; no regression past defined thresholds

15. The Shifting Model / Harness Boundary




Figure 15. The model upgrade paradox.

The line between model and harness is not fixed. As models receive more agentic training — calling tools, planning, recovering from errors natively — work that once had to be engineered into the harness increasingly lives in the weights. This moves the harness's frontier rather than erasing it: undifferentiated work (basic looping, retry logic) migrates inward, while the functions that are external *by nature* — identity, delegation, provenance, isolation, auditability, cost control, integration — remain harness territory. As models are trusted with more consequential authority, those control functions become more important, not less. Treat specific architectural recommendations as provisional against this moving boundary.

The model upgrade paradox. Today's measured attack outcomes partly reflect security by incompetence: WASP found that agents begin executing adversarial instructions 16–86% of the time but complete the attacker's full goal only 0–17% of the time — the gap exists not because a defense fired, but because the model wasn't capable enough to finish the job. That protection is borrowed, and a model upgrade repays it: swapping in a more capable model, with the harness unchanged, converts previously failed exploits into completed ones, because better planning and execution serve injected goals as readily as legitimate ones. The operational rule: **treat every model upgrade as a security-relevant change** — re-run the adversarial suite (§14) against the exact production configuration before rollout, and assume that any partial-compromise rate observed under the old model is a floor, not a ceiling, under the new one.

16. Limitations & Open Problems




Figure 16. Limitations and open problems.

- **Prompt injection is not solved.** No reliable defense exists against a model following instructions embedded in untrusted content it reads. Everything in §6.3–§6.4 *mitigates and contains*; nothing eliminates. This remains the central open problem for agents that read the open web.

- **Provenance-aware policy is young.** Taint tracking through summarization, trusted extractors, and provenance-gated authorization are the right direction but lack mature, standardized implementations. The same holds for **mandate-level (trajectory) authorization**: continuously validating an agent's structural behavior against the original user intent (§6.4.1) is the correct response to laundering-style attacks, but no standardized mechanism for intent digests or risk-tier drift detection exists yet.
- **Long-horizon reliability degrades.** If ten independent *required* steps each succeed with probability 0.95, the end-to-end success probability is $0.95^{10} \approx 60\%$. Real agent errors are often correlated, so treat this as an illustration rather than an estimator — but the qualitative point stands: robust behavior over hundreds of steps remains hard.
- **Context management is heuristic.** What to keep, compress, or retrieve is hand-tuned; retrieval quality is often the hidden bottleneck.
- **Agent evaluation — especially adversarial evaluation — is immature.** Scores are harness-relative (§8); adaptive attacks dramatically outperform fixed sets; and reproducibility is limited to forensic reconstruction (§6.6).
- **Agent identity standards are in flight.** NIST's concept work on software and AI-agent identity and authorization is early; today's delegation and capability mechanisms are bespoke.
- **The model/harness boundary moves (§15),** so today's division of responsibilities is provisional.

These are not reasons to avoid building; they are where the engineering effort is currently concentrated.

17. Harness vs. Model: A Clarifying Summary

Defining the Trust Boundary: Model vs. Harness		
	Model	Harness
Role	Proposes actions, generates text.	Authorizes, executes, verifies, records.
State	No durable, trusted state.	Owns memory, authority and the audit trail.
Capabilities	Reasoning, language.	Tools, context, identity, execution control.
Trust Position	Inside the blast radius.	The perimeter and the record.
Determines	The ceiling of what is possible.	How much is realized – and how safely.

Figure 17. Harness vs. model — a clarifying summary of roles and trust position.

	Model	Harness
Role	Proposes actions, generates text	Authorizes, executes, verifies, records
State	No durable <i>trusted</i> state	Owns memory, authority, and audit
Provides	Reasoning, language	Tools, context, identity, control
Determines	The ceiling of what's possible	How much is realized — and how safely
Trust position	Inside the blast radius	The perimeter and the record
Analogy	The horse (raw power)	Collar, traces, reins (direction and control)

A frontier model without a good harness is a powerful engine with no drivetrain — capability that never reaches the road.

18. Closing Note




Figure 18. *The model is the engine; the harness is the difference.*

The harness is where most of the *engineering* of an AI system happens. Within it, four disciplines carry the weight: **the tools that connect the model to the world; the context and memory that decide what it sees and retains; the authority and execution controls that decide what it may actually do; and the evaluation that proves any of it works — including against an adversary.** Model quality sets the ceiling; harness quality determines how much of it is realized and whether the realized power remains under the control of the person it serves. As models grow more capable and are delegated more authority, that work becomes more central to the field, not less.

19. References & Further Reading

Reasoning and acting

- Yao et al. (2022). *ReAct: Synergizing Reasoning and Acting in Language Models*.
- Wei et al. (2022). *Chain-of-Thought Prompting Elicits Reasoning in Large Language Models*.
- Shinn et al. (2023). *Reflexion: Language Agents with Verbal Reinforcement Learning*.

Tool use

- Schick et al. (2023). *Toolformer: Language Models Can Teach Themselves to Use Tools*.

Evaluation harnesses & benchmarks

- Gao et al. *A Framework for Few-Shot Language Model Evaluation* ([lm-evaluation-harness](#), EleutherAI).
- Jiménez et al. (2023). *SWE-bench: Can Language Models Resolve Real-World GitHub Issues?*

Security (key sources verified July 2026)

- Hardy, N. (1988). *The Confused Deputy (or why capabilities might have been invented)*.
- Greshake et al. (2023). *Not What You've Signed Up For: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection*.
- US AISI / CAISI at NIST (Jan 2025). *Technical Blog: Strengthening AI Agent Hijacking Evaluations — AgentDojo-based; source of the adaptive-attack (11%→81%) and repeated-attempt findings, and of the insight that evaluations must be adaptive and multi-attempt*. [nist.gov/news-events/news/2025/01/technical-blog-strengthening-ai-agent-hijacking-evaluations](#)
- Evtimov, Zharmagambetov, Grattafiori, Guo, Chaudhuri (2025). *WASP: Benchmarking Web Agent Security Against Prompt Injection Attacks*. arXiv:2504.18575 — source of the partial-vs-complete compromise gap (16–86% begin executing injected instructions; 0–17% complete the attacker goal).
- OWASP GenAI Security Project (2026). *Top 10 for Agentic Applications* — the ASI taxonomy referenced in §6.2, §6.3, and §11, including ASIO4 (agentic supply chain), ASIO6 (memory & context poisoning), ASIO7 (insecure inter-agent communication), ASIO8 (cascading failures), and ASIO10 (rogue agents).
- NIST CSRC (2026, initial public draft). *Accelerating the Adoption of Software and AI Agent Identity and Authorization* (concept paper).

Identity & authorization standards

- IETF RFC 8693. *OAuth 2.0 Token Exchange* — downstream, down-scoped delegation.
- SPIFFE / SPIRE — workload identity and attestation frameworks.
- NIST NGAC — attribute-based, value-aware access-control modeling.

Practice

- Provider documentation on tool use/function calling, agent frameworks, and evaluation for current, versioned APIs.

(The security sources above were verified against live primary sources in July 2026; the reasoning/tool-use references are standard works cited by author–year — consult primary sources for exact venues. One precision note: the widely quoted "57%→80% over 25 attempts" repetition figure attributed to the NIST blog circulates in secondary summaries; the primary source confirms the multi-attempt insight and the 11%→81% adaptation finding, so this document leans on those.)